

```

1 #Importing all the required libraries
2
3 import os
4 import pandas as pd
5 import numpy as np
6 import matplotlib.pyplot as plt
7 from scipy.stats import chi2_contingency
8 import seaborn as sns
9 from random import randrange, uniform

```

```

1 #Reading the train data
2
3 train_dataset = pd.read_csv("/content/train_Df64byy.csv")
4 train_dataset.head()

```

	ID	City_Code	Region_Code	Accommodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health Indicator	Holc
0	1	C3	3213	Rented	Individual	36	36	No	X1	
1	2	C5	1117	Owned	Joint	75	22	No	X2	
2	3	C5	3732	Owned	Individual	32	32	No	NaN	
3	4	C24	4378	Owned	Joint	52	48	No	X1	
4	5	C8	2190	Rented	Individual	44	44	No	X2	

```

1 #Describing the train data
2
3 train_dataset.describe()

```

	ID	Region_Code	Upper_Age	Lower_Age	Holding_Policy_Type	Reco_Policy_Cat	Reco_Policy_Premium
count	66720.000000	66720.000000	66720.000000	66720.000000	40123.000000	66720.000000	66720.000000
mean	29600.349460	1733.354332	44.846927	42.734937	2.442664	15.112125	14183.940833
std	15003.069766	1424.021443	17.299650	17.310907	1.025191	6.340442	6585.290312
min	1.000000	1.000000	18.000000	16.000000	1.000000	1.000000	2280.000000
25%	16679.750000	527.000000	28.000000	27.000000	1.000000	12.000000	9252.000000

```

1 #Returning the data types of train
2
3 train_dataset.dtypes

```

```

ID                int64
City_Code         object
Region_Code       float64
Accomodation_Type object
Reco_Insurance_Type object
Upper_Age         float64
Lower_Age         float64
Is_Spouse         object
Health Indicator  object
Holding_Policy_Duration object
Holding_Policy_Type float64
Reco_Policy_Cat   float64
Reco_Policy_Premium float64
Response          float64
dtype: object

```

```

1 #Returning the no of classified and non classified
2
3 print("No of observations classified as 1 are : ", len(train_dataset[train_dataset['Response']==1]))
4
5 print("No of observations classified as 0 are : ", len(train_dataset[train_dataset['Response']==0]))

```

```

No of observations classified as 1 are : 16080
No of observations classified as 0 are : 50640

```

```

1 #Checking no of missing values in train data
2
3 missing_val = pd.DataFrame(train_dataset.isnull().sum())
4 missing_val

```

	0
ID	0
City_Code	1
Region_Code	1
Accomodation_Type	1
Reco_Insurance_Type	1
Upper_Age	1
Lower_Age	1
Is_Spouse	1
Health Indicator	11692
Holding_Policy_Duration	20252
Holding_Policy_Type	20252
Reco_Policy_Cat	1
Reco_Policy_Premium	1
Response	1

```

1 #Returning the value counts for each unique value in Holding_Policy_Duration Column
2
3 train_dataset['Holding_Policy_Duration'].value_counts()

```

```

1.0    5889
14+    5647
2.0    5582
3.0    4724
4.0    3612

```

5.0	3103
6.0	2495
7.0	2161
8.0	1726
9.0	1451
10.0	1070
11.0	709
13.0	688
12.0	659
14.0	607

Name: Holding_Policy_Duration, dtype: int64

```
1 #Replacing 14+ value in Holding_Policy_Duration with 15
2
3 # mapping = {'14+':'15'}
4 train_dataset['Holding_Policy_Duration'] = train_dataset['Holding_Policy_Duration'].replace("14+", "15")
```

```
1 #Returning the value counts for each unique value in Health_Indicator Column
2
3 train_dataset["Health Indicator"].value_counts()
```

X1	17087
X2	13522
X3	8890
X4	7550
X5	2242
X6	1677
X7	258
X8	105
X9	86

Name: Health Indicator, dtype: int64

```
1 #Returning the value counts for each unique value in Holding_Policy_Type Column
2
3 train_dataset["Holding_Policy_Type"].value_counts()
```

3.0	17432
1.0	10647
2.0	6556

4.0 5488

Name: Holding_Policy_Type, dtype: int64

```

1 #Checking the percentage of missing values for train data
2
3 missing_val = missing_val.reset_index()
4 missing_val = missing_val.rename(columns = {'index':'Variables', 0: 'Missing_percentage'})
5 missing_val['Missing_percentage'] = (missing_val['Missing_percentage']/len(train_dataset))*100
6 missing_val = missing_val.sort_values('Missing_percentage', ascending = False).reset_index(drop = True)
7 missing_val

```

	Variables	Variables	Missing_percentage
0	1	City_Code	0.000003
1	2	Region_Code	0.000003
2	3	Accommodation_Type	0.000003
3	4	Reco_Insurance_Type	0.000003
4	5	Upper_Age	0.000003
5	6	Lower_Age	0.000003
6	7	Is_Spouse	0.000003
7	11	Reco_Policy_Cat	0.000003
8	12	Reco_Policy_Premium	0.000003
9	13	Response	0.000003
10	0	ID	0.000000
11	8	Health Indicator	0.000000
12	9	Holding_Policy_Duration	0.000000
13	10	Holding Policy Type	0.000000

```

1 #Imputing the missing values for required columns
2

```

```
2
3 from sklearn.impute import SimpleImputer
4
5 impute_size = SimpleImputer(strategy = "most_frequent")
6 train_dataset['Health Indicator'] = impute_size.fit_transform(train_dataset[['Health Indicator']])
7 train_dataset['Holding_Policy_Duration'] = impute_size.fit_transform(train_dataset[['Holding_Policy_Duration']])
8 train_dataset['Holding_Policy_Type'] = train_dataset['Holding_Policy_Type'].fillna(train_dataset['Holding_Policy_Type'])

1 #Checking the percentage of missing values for train data after imputing the missing values
2
3 missing_val = pd.DataFrame(train_dataset.isnull().sum())
4 missing_val = missing_val.reset_index()
5 missing_val = missing_val.rename(columns = {'index':'Variables', 0: 'Missing_percentage'})
6 missing_val['Missing_percentage'] = (missing_val['Missing_percentage']/len(train_dataset))*100
7 missing_val = missing_val.sort_values('Missing_percentage', ascending = False).reset_index(drop = True)
8 missing_val
```

Variables Missing_percentage

```

1 #Still we can some missing values percentage, so we have to drop columns which have nan value
2
3 train_dataset.dropna(how='any')

```

	ID	City_Code	Region_Code	Accommodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator
0	1	C3	3213.0	Rented	Individual	36.0	36.0	No	>
1	2	C5	1117.0	Owned	Joint	75.0	22.0	No	>
2	3	C5	3732.0	Owned	Individual	32.0	32.0	No	>
3	4	C24	4378.0	Owned	Joint	52.0	48.0	No	>
4	5	C8	2190.0	Rented	Individual	44.0	44.0	No	>
...
50877	35040	C1	102.0	Rented	Joint	48.0	47.0	Yes	>
50878	35041	C1	332.0	Owned	Individual	70.0	70.0	No	>
50879	35042	C6	1165.0	Owned	Joint	57.0	56.0	Yes	>
50880	35043	C11	1032.0	Owned	Joint	72.0	69.0	Yes	>
50881	35044	C1	729.0	Rented	Individual	25.0	25.0	No	>

50882 rows × 14 columns

```

1 #Checking the percentage of missing values for train data after dropping the missing values
2
3 missing_val = pd.DataFrame(train_dataset.isnull().sum())
4 missing_val = missing_val.reset_index()
5 missing_val = missing_val.rename(columns = {'index':'Variables', 0: 'Missing_percentage'})
6 missing_val['Missing_percentage'] = (missing_val['Missing_percentage']/len(train_dataset))*100
7 missing_val1 = missing_val.sort_values('Missing_percentage', ascending = False).reset_index(drop = True)
8 missing_val1

```

	Variables	Missing_percentage
0	ID	0.0
1	City_Code	0.0
2	Region_Code	0.0
3	Accommodation_Type	0.0
4	Reco_Insurance_Type	0.0
5	Upper_Age	0.0
6	Lower_Age	0.0
7	Is_Spouse	0.0
8	Health Indicator	0.0
9	Holding_Policy_Duration	0.0
10	Holding_Policy_Type	0.0
11	Reco_Policy_Cat	0.0
12	Reco_Policy_Premium	0.0
13	Response	0.0

I'm trying to consider only 'Reco_Policy_Premium' column for outliers because it's the only column which deals with money and all other columns belongs to ID, Age Duration Etc

```

1 #Process to detect outliers with the help of quartile
2
3 max_thresold = train_dataset['Reco_Policy_Premium'].quantile(0.95)
4 max_thresold

```

26852.0

```

1 #Checking the data which are outliers with max_thresold

```



```

2
3 train_dataset[train_dataset['Reco_Policy_Premium'] > max_threshold]

```

	ID	City_Code	Region_Code	Accomodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator
1	2	C5	1117.0	Owned	Joint	75.0	22.0	No	>
7	8	C1	3175.0	Owned	Joint	75.0	73.0	Yes	>
8	9	C15	3497.0	Owned	Joint	52.0	43.0	No	>
48	49	C2	2858.0	Owned	Joint	57.0	55.0	Yes	>
49	50	C1	85.0	Owned	Joint	73.0	68.0	Yes	>
...
50793	34956	C1	3014.0	Owned	Joint	47.0	46.0	Yes	>
50798	34961	C2	484.0	Owned	Joint	69.0	68.0	Yes	>
50812	34975	C1	907.0	Owned	Joint	75.0	72.0	Yes	>
50842	35005	C1	4370.0	Owned	Joint	71.0	68.0	No	>
50852	35015	C1	1383.0	Owned	Joint	59.0	58.0	Yes	>

2544 rows × 14 columns

```

1 #Process to detect outliers with the help of quartile
2
3 min_threshold = train_dataset['Reco_Policy_Premium'].quantile(0.05)
4 min_threshold

```

5224.0

```

1 #Checking the data which are outliers with min_threshold
2
3 train_dataset[train_dataset['Reco_Policy_Premium'] < min_threshold]

```

	ID	City_Code	Region_Code	Accomodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator
10	11	C28	600.0	Owned	Individual	21.0	21.0	No	>
15	16	C3	1484.0	Rented	Individual	20.0	20.0	No	>
22	23	C25	787.0	Rented	Individual	18.0	18.0	No	>
27	28	C9	855.0	Rented	Individual	21.0	21.0	No	>
46	47	C3	1475.0	Rented	Individual	21.0	21.0	No	>
...	
50775	34938	C3	1246.0	Owned	Individual	19.0	19.0	No	>
50814	34977	C6	155.0	Rented	Individual	28.0	28.0	No	>
50822	34985	C1	1984.0	Rented	Individual	19.0	19.0	No	>
50836	34999	C3	1418.0	Rented	Individual	20.0	20.0	No	>
50874	35037	C27	3277.0	Owned	Individual	29.0	29.0	No	>

2537 rows × 14 columns

```
1 train_dataset = train_dataset[(train_dataset['Reco_Policy_Premium'] < max_threshold) & (train_dataset['Reco_Policy_Premium'] > min_threshold)]
```

```
1 train_dataset
```

	ID	City_Code	Region_Code	Accomodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator
0	1	C3	3213.0	Rented	Individual	36.0	36.0	No	>
2	3	C5	3732.0	Owned	Individual	32.0	32.0	No	>
3	4	C24	4378.0	Owned	Joint	52.0	48.0	No	>
4	5	C8	2190.0	Rented	Individual	44.0	44.0	No	>
5	6	C9	1785.0	Rented	Individual	52.0	52.0	No	>
...

```

1 #Returning the no of classified and non classified after imputing missing values and removing outliers from train data
2
3 print("No of observations classified as 1 are : ", len(train_dataset[train_dataset['Response'] == 1]))
4
5 print("No of observations classified as 0 are : ", len(train_dataset[train_dataset['Response'] == 0]))

No of observations classified as 1 are : 16080
No of observations classified as 0 are : 50640

15770 rows x 10 columns

1 #Describing the test_data
2 test_dataset =pd.read_csv('/content/test_YCcrUnU.csv')
3 test_dataset.describe()

```

```

      ID   Region_Code   Upper_Age   Lower_Age   Holding_Policy_Type   Reco_Policy_Cat   Reco_Policy_Premium
1  #Checking the data types for test data
2  test_dataset = pd.read_csv("/content/test_YCcrUnU.csv")
3
4  test_dataset.dtypes

ID                int64
City_Code         object
Region_Code       int64
Accomodation_Type object
Reco_Insurance_Type object
Upper_Age         int64
Lower_Age         int64
Is_Spouse         object
Health Indicator  object
Holding_Policy_Duration object
Holding_Policy_Type float64
Reco_Policy_Cat   int64
Reco_Policy_Premium float64
dtype: object

1  #Checking the percentage of missing values for test data
2
3  missing_val = pd.DataFrame(test_dataset.isnull().sum())
4  missing_val = missing_val.reset_index()
5  missing_val = missing_val.rename(columns = {'index':'Variables', 0: 'Missing_percentage'})
6  missing_val['Missing_percentage'] = (missing_val['Missing_percentage']/len(test_dataset))*100
7  missing_val = missing_val.sort_values('Missing_percentage', ascending = False).reset_index(drop = True)
8  missing_val

```

	Variables	Missing_percentage
0	Holding_Policy_Duration	39.454254
1	Holding_Policy_Type	39.454254
2	Health_Indicator	23.054345
3	ID	0.000000
4	City_Code	0.000000
5	Region_Code	0.000000
6	Accommodation_Type	0.000000
7	Reco_Insurance_Type	0.000000

```

1 #Returning the value counts for each unique value in Holding_Policy_Duration Column
2
3 test_dataset['Holding_Policy_Duration'].value_counts()

```

```

14+      1892
1.0      1891
2.0      1772
3.0      1606
4.0      1205
5.0       992
6.0       903
7.0       664
8.0       569
9.0       493
10.0      333
11.0      254
13.0      221
14.0      211
12.0      196

```

```
Name: Holding_Policy_Duration, dtype: int64
```

```

1 #Replacing 14+ value with 15 for 'Holding_Policy_Duration' column
2
3 test_dataset['Holding_Policy_Duration'] = test_dataset['Holding_Policy_Duration'].replace("14+", "15")

```

```

1 #Returning the value counts for each unique value in Health Indicator Column
2
3 test_dataset['Health Indicator'].value_counts()

```

```

X1      5614
X2      4516
X3      2846
X4      2442
X5       681
X6       514
X7        96
X8        41
X9        28

```

Name: Health Indicator, dtype: int64

```

1 #Returning the value counts for each unique value in Holding_Policy_Type Column
2
3 test_dataset['Holding_Policy_Type'].value_counts()

```

```

3.0      5572
1.0      3574
2.0      2150
4.0      1906

```

Name: Holding_Policy_Type, dtype: int64

```

1 #Imputing missing values for test data
2
3 impute_size = SimpleImputer(strategy = "most_frequent")
4 test_dataset['Health Indicator'] = impute_size.fit_transform(test_dataset[['Health Indicator']])
5 test_dataset['Holding_Policy_Duration'] = impute_size.fit_transform(test_dataset[['Holding_Policy_Duration']])
6 test_dataset['Holding_Policy_Type'] = test_dataset['Holding_Policy_Type'].fillna(test_dataset['Holding_Policy_Type'].me

```

```

1 #Checking the percentage of missing values for test data after imputing the missing values
2

```

```

3 missing_val = pd.DataFrame(test_dataset.isnull().sum())
4 missing_val = missing_val.reset_index()
5 missing_val = missing_val.rename(columns = {'index':'Variables', 0: 'Missing_percentage'})
6 missing_val['Missing_percentage'] = (missing_val['Missing_percentage']/len(test_dataset))*100

```

```

6 missing_val[missing_percentage] = (missing_val[missing_percentage] / len(test_dataset)) * 100
7 missing_val = missing_val.sort_values('Missing_percentage', ascending = False).reset_index(drop = True)
8 missing_val

```

	Variables	Missing_percentage
0	ID	0.0
1	City_Code	0.0
2	Region_Code	0.0
3	Accommodation_Type	0.0
4	Reco_Insurance_Type	0.0
5	Upper_Age	0.0
6	Lower_Age	0.0
7	Is_Spouse	0.0
8	Health Indicator	0.0
9	Holding_Policy_Duration	0.0
10	Holding_Policy_Type	0.0
11	Reco_Policy_Cat	0.0
12	Reco Policy Premium	0.0

```

1 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
2
3 le = LabelEncoder()
4 objList = train_dataset.select_dtypes(include="object").columns
5
6 for feat in objList:
7     train_dataset[feat] = le.fit_transform(train_dataset[feat].astype(str))

```

```

1 train_dataset

```

	ID	City_Code	Region_Code	Accomodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator
0	1	22	3213	1	0	36	36	0	
1	2	31	1117	0	1	75	22	0	
2	3	31	3732	0	0	32	32	0	
3	4	16	4378	0	1	52	48	0	
4	5	34	2190	1	0	44	44	0	
...
66715	50878	30	845	1	0	22	22	0	
66716	50879	31	4188	1	0	27	27	0	
66717	50880	0	442	1	0	63	63	0	
66718	50881	0	4	0	1	71	49	0	
66719	50882	22	3866	1	0	24	24	0	

66720 rows × 14 columns

```

1 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
2
3 le = LabelEncoder()
4 objList = test_dataset.select_dtypes(include="object").columns
5
6 for feat in objList:
7     test_dataset[feat] = le.fit_transform(test_dataset[feat].astype(str))

1 test_dataset.head()

```


	ID	City_Code	Region_Code	Accommodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator	
0	50883	0	156	0	0	30	30	0	0	
1	50884	30	7	0	1	69	68	1	0	

```
1 X = train_dataset.drop("Response", axis=1) #Feature Matrix
2 y = train_dataset["Response"]
```

4	50887	0	951	0	0	75	75	0	2	
---	-------	---	-----	---	---	----	----	---	---	--

```
1 X.head()
```

	ID	City_Code	Region_Code	Accommodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health_Indicator	Holds_Out
0	1	22	3213	1	0	36	36	0	0	
2	3	31	3732	0	0	32	32	0	0	
3	4	16	4378	0	1	52	48	0	0	
4	5	34	2190	1	0	44	44	0	1	
5	6	35	1785	1	0	52	52	0	1	

```
1 # separate dataset into train and test
2 from sklearn.model_selection import train_test_split
3 X_train, X_val, y_train, y_val = train_test_split(
4     X,
5     y,
6     test_size=0.2,
7     random_state=39)
8
9 X_train.shape, X_val.shape

((53376, 13), (13344, 13))
```

▼ Handling Imablance

```
1  from sklearn.datasets import make_classification
2  from sklearn.linear_model import LogisticRegression
3  from sklearn.dummy import DummyClassifier
4  from sklearn.model_selection import train_test_split
5  from sklearn.metrics import roc_curve
6  from sklearn.metrics import roc_auc_score
7  from matplotlib import pyplot
8
9  # plot no skill and model roc curves
10 def plot_roc_curve(test_y, naive_probs, model_probs):
11     # plot naive skill roc curve
12     fpr, tpr, _ = roc_curve(test_y, naive_probs)
13     pyplot.plot(fpr, tpr, linestyle='--', label='No Skill')
14     # plot model roc curve
15     fpr, tpr, _ = roc_curve(test_y, model_probs)
16     pyplot.plot(fpr, tpr, marker='.', label='Logistic')
17     # axis labels
18     pyplot.xlabel('False Positive Rate')
19     pyplot.ylabel('True Positive Rate')
20     # show the legend
21     pyplot.legend()
22     # show the plot
23     pyplot.show()
24
25 from sklearn.metrics import precision_recall_curve
26
27 # no skill model, stratified random class predictions
28 model = DummyClassifier(strategy='stratified')
29 model.fit(X_train,y_train)
30 yhat = model.predict_proba(X_val)
31 pos_probs = yhat[:, 1]
32 # calculate the precision-recall auc
33 precision, recall, _ = precision_recall_curve(y_val, pos_probs)
34 auc_score = auc(recall, precision)
35 print('No Skill PR AUC: %.3f' % auc_score)
```

No Skill PR AUC: 0.333

```

1 # skilled model
2 model = LogisticRegression(solver='lbfgs')
3 model.fit(X_train, y_train)
4 yhat = model.predict_proba(X_val)
5 model_probs = yhat[:, 1]
6 # calculate roc auc
7 roc_auc = roc_auc_score(y_val, model_probs)
8 print('Logistic ROC AUC %.3f' % roc_auc)
9

```

Logistic ROC AUC 0.501

```

1 # example of evaluating a decision tree with random oversampling
2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import cross_val_score
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from sklearn.tree import DecisionTreeClassifier
7 from imblearn.pipeline import Pipeline
8 from imblearn.over_sampling import RandomOverSampler

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/externals/six.py:31: FutureWarning: The module is deprecated in version
"(https://pypi.org/project/six/).", FutureWarning)
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The sklearn.neighbors.base module
warnings.warn(message, FutureWarning)

```

```

1 # define pipeline
2 steps = [('over', RandomOverSampler()), ('model', DecisionTreeClassifier())]
3 pipeline = Pipeline(steps=steps)
4 # evaluate pipeline
5 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
6 scores = cross_val_score(pipeline, X, y, scoring='f1_micro', cv=cv, n_jobs=-1)
7 score = mean(scores)

```

```
8 print('F-measure: %.3f' % score)
```

```
F-measure: 0.668
```

```
1 from numpy import where
2 from collections import Counter
3 from imblearn.over_sampling import SMOTE
4 # summarize class distribution
5 counter = Counter(y)
6 print(counter)
7 # scatter plot of examples by class label
8 # transform the dataset
9 oversample = SMOTE()
10 X, y = oversample.fit_resample(X, y)
```

```
Counter({0: 34779, 1: 10991})
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated
warnings.warn(msg, category=FutureWarning)
```

```
1 # summarize the new class distribution
2 counter = Counter(y)
3 print(counter)
```

```
Counter({0: 34779, 1: 34779})
```

```
1 # define model
2 model = DecisionTreeClassifier()
3 # evaluate pipeline
4 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
5 scores = cross_val_score(model, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
6 print('Mean ROC AUC: %.3f' % mean(scores))
```

```
Mean ROC AUC: 0.758
```

```
1 # define pipeline
2 steps = [('over', SMOTE()), ('model', DecisionTreeClassifier())]
```

```

3 pipeline = Pipeline(steps=steps)
4 # evaluate pipeline
5 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
6 scores = cross_val_score(pipeline, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
7 print('Mean ROC AUC: %.3f' % mean(scores))

```

Mean ROC AUC: 0.758

```

1 from imblearn.over_sampling import BorderlineSMOTE
2 # summarize class distribution
3 counter = Counter(y)
4 print(counter)
5 # transform the dataset
6 oversample = BorderlineSMOTE()
7 X, y = oversample.fit_resample(X_train, y_train)
8 # summarize the new class distribution
9 counter = Counter(y)
10 # summarize the new class distribution
11 counter = Counter(y)
12 print(counter)

```

Counter({0: 9900, 1: 100})

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated
warnings.warn(msg, category=FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated
warnings.warn(msg, category=FutureWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function safe_indexing is deprecated
warnings.warn(msg, category=FutureWarning)

Counter({1: 27835, 0: 27835})

```

1 # combined SMOTE and Tomek Links sampling for imbalanced classification
2 from numpy import mean
3 from sklearn.datasets import make_classification
4 from sklearn.model_selection import cross_val_score
5 from sklearn.model_selection import RepeatedStratifiedKFold
6 from imblearn.pipeline import Pipeline
7 from sklearn.tree import DecisionTreeClassifier
8 from imblearn.combine import SMOTETomek

```

```

9  from imblearn.under_sampling import TomekLinks
10
11  # define model
12  model = DecisionTreeClassifier()
13  # define sampling
14  resample = SMOTETomek(tomek=TomekLinks(sampling_strategy='majority'))
15  # define pipeline
16  pipeline = Pipeline(steps=[('r', resample), ('m', model)])
17  # define evaluation procedure
18  cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
19  # evaluate model
20  scores = cross_val_score(pipeline, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
21  # summarize performance
22  print('Mean ROC AUC: %.3f' % mean(scores))

```

Mean ROC AUC: 0.529

```

1  # fit a logistic regression model on an imbalanced classification dataset
2  from numpy import mean
3  from sklearn.datasets import make_classification
4  from sklearn.model_selection import cross_val_score
5  from sklearn.model_selection import RepeatedStratifiedKFold
6  from sklearn.linear_model import LogisticRegression
7
8  # define model
9  weights = {0:0.01, 1:1.0}
10 model = LogisticRegression(solver='lbfgs', class_weight=weights)
11 model = LogisticRegression(solver='lbfgs')
12 # define evaluation procedure
13 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
14 # evaluate model
15 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
16 # summarize performance
17 print('Mean ROC AUC: %.3f' % mean(scores))

```

Mean ROC AUC: 0.507

```

1  # calculate heuristic class weighting

```

```

2 from sklearn.utils.class_weight import compute_class_weight
3 # calculate class weighting
4 weighting = compute_class_weight('balanced', [0,1], y_train)
5 print(weighting)

```

```
[0.65773307 2.08495616]
```

```

1 # define model
2 model = LogisticRegression(solver='lbfgs', class_weight='balanced')
3 # define evaluation procedure
4 cv = RepeatedStratifiedKFold(n_splits=25, n_repeats=10, random_state=1)
5 # evaluate model
6 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
7 # summarize performance
8 print('Mean ROC AUC: %.3f' % mean(scores))

```

```
Mean ROC AUC: 0.552
```

```

1 # define model
2 model = LogisticRegression(solver='lbfgs')
3 # define grid
4 balance = [{0:100,1:1}, {0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}]
5 param_grid = dict(class_weight=balance)
6 # define evaluation procedure
7 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
8 # define grid search
9 grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=cv,
10 scoring='roc_auc')
11 # execute the grid search
12 grid_result = grid.fit(X_train, y_train)
13 # report the best configuration
14 print('Best: %f using %s' % (grid_result.best_score_, grid_result.best_params_))
15 # report all configurations
16 means = grid_result.cv_results_['mean_test_score']
17 stds = grid_result.cv_results_['std_test_score']
18 params = grid_result.cv_results_['params']
19 for mean, stdev, param in zip(means, stds, params):
20     print('%f (%f) with: %r' % (mean, stdev, param))

```

```

Best: 0.561343 using {'class_weight': {0: 1, 1: 10}}
0.478032 (0.009518) with: {'class_weight': {0: 100, 1: 1}}
0.479372 (0.011027) with: {'class_weight': {0: 10, 1: 1}}
0.506829 (0.009842) with: {'class_weight': {0: 1, 1: 1}}
0.561343 (0.011775) with: {'class_weight': {0: 1, 1: 10}}
0.555669 (0.011346) with: {'class_weight': {0: 1, 1: 100}}

```

```

1 # fit a decision tree on an imbalanced classification dataset
2 from numpy import mean
3 # define model
4 model = DecisionTreeClassifier()
5 # define evaluation procedure
6 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
7 # evaluate model
8 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
9 # summarize performance
10 print('Mean ROC AUC: %.3f' % mean(scores))

```

Mean ROC AUC: 0.544

```

1 # define model
2 model = DecisionTreeClassifier(class_weight='balanced')
3 # define evaluation procedure
4 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
5 # evaluate model
6 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
7 # summarize performance
8 print('Mean ROC AUC: %.3f' % mean(scores))

```

Mean ROC AUC: 0.542

```

1 #define model
2 model = DecisionTreeClassifier()
3 # define grid
4 balance = [{0:100,1:1}, {0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}]
5 param_grid = dict(class_weight=balance)
6 # define evaluation procedure

```



```

3 # define evaluation procedure
7 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
8 # define grid search
9 grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=cv,
10 scoring='roc_auc')
11 # execute the grid search
12 grid_result = grid.fit(X_train, y_train)
13 # report the best configuration
14 print('Best: %f using %s' % (grid_result.best_score_, grid_result.best_params_))
15 # report all configurations
16 means = grid_result.cv_results_['mean_test_score']
17 stds = grid_result.cv_results_['std_test_score']
18 params = grid_result.cv_results_['params']
19 for mean, stdev, param in zip(means, stds, params):
20     print('%f (%f) with: %r' % (mean, stdev, param))

```

```

Best: 0.544592 using {'class_weight': {0: 10, 1: 1}}
0.542249 (0.008030) with: {'class_weight': {0: 100, 1: 1}}
0.544592 (0.011834) with: {'class_weight': {0: 10, 1: 1}}
0.543369 (0.009397) with: {'class_weight': {0: 1, 1: 1}}
0.536144 (0.008191) with: {'class_weight': {0: 1, 1: 10}}
0.536047 (0.006741) with: {'class_weight': {0: 1, 1: 100}}

```

```

1 # define model
2 model = SVC(gamma='scale', class_weight='balanced')
3 # define evaluation procedure
4 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
5 # evaluate model
6 scores = cross_val_score(model, X, y, scoring='roc_auc', cv=cv, n_jobs=-1)
7 # summarize performance
8 print('Mean ROC AUC: %.3f' % mean(scores))

```

```

1 from sklearn.metrics import roc_auc_score
2 from keras.layers import Dense
3 from keras.models import Sequential
4 # define the neural network model
5 def define_model(n_input):
6     # define model
7     model = Sequential()

```

```
8 # define first hidden layer and visible layer
9 model.add(Dense(10, input_dim=n_input, activation='relu',
10 kernel_initializer='he_uniform'))
11 # define output layer
12 model.add(Dense(1, activation='sigmoid'))
13 # define loss and optimizer
14 model.compile(loss='binary_crossentropy', optimizer='sgd')
15 return model
```

```
1 # define the model
2 n_input = X_train.shape[1]
3 model = define_model(n_input)
```

```
1 # fit model
2 model.fit(X_train,y_train,epochs=100, verbose=1)
3 # make predictions on the test dataset
4 yhat = model.predict(y_val)
5 # evaluate the ROC AUC of the predictions
6 score = roc_auc_score(testy, yhat)
7 print('ROC AUC: %.3f' % score)
```

```
1 model = define_model(n_input)
2 # fit model
3 weights = {0:1, 1:100}
4 history = model.fit(X_train,y_train, class_weight=weights, epochs=100, verbose=1)
5 # evaluate model
6 yhat = model.predict(X_val)
7 score = roc_auc_score(y_val, yhat)
8 print('ROC AUC: %.3f' % score)
```

```
Epoch 1/100
1145/1145 [=====] - 2s 2ms/step - loss: 12005825.8617
Epoch 2/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3963
Epoch 3/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3979
Epoch 4/100
```

```
1145/1145 [=====] - 2s 2ms/step - loss: 3.3947
Epoch 5/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3975
Epoch 6/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3968
Epoch 7/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3958
Epoch 8/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3965
Epoch 9/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3964
Epoch 10/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3955
Epoch 11/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3959
Epoch 12/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3961
Epoch 13/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3973
Epoch 14/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3957
Epoch 15/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3970
Epoch 16/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3957
Epoch 17/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3955
Epoch 18/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3959
Epoch 19/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3939
Epoch 20/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3957
Epoch 21/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3974
Epoch 22/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3978
Epoch 23/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3959
Epoch 24/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3956
Epoch 25/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3973
```

```
Epoch 26/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3957
Epoch 27/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3961
Epoch 28/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3959
Epoch 29/100
1145/1145 [=====] - 2s 2ms/step - loss: 3.3967
- ... -
```

```
1 # fit xgboost on an imbalanced classification dataset
2 from numpy import mean
3 from sklearn.model_selection import cross_val_score
4 from sklearn.model_selection import RepeatedStratifiedKFold
5 from xgboost import XGBClassifier
```

```
1 # define model
2 model = XGBClassifier()
3 # define evaluation procedure
4 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
5 # evaluate model
6 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
7 # summarize performance
8 print('Mean ROC AUC: %.5f' % mean(scores))
```

Mean ROC AUC: 0.62498

```
1 # define model
2 model = XGBClassifier(scale_pos_weight=99)
3 # define evaluation procedure
4 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
5 # evaluate model
6 scores = cross_val_score(model, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
7 # summarize performance
8 print('Mean ROC AUC: %.5f' % mean(scores))
```

Mean ROC AUC: 0.62435

```
1 # define grid
```

```

1  # define grid
2  weights = [1, 10, 25, 50, 75, 99, 100, 1000]
3  param_grid = dict(scale_pos_weight=weights)
4  # define evaluation procedure
5  cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
6  # define grid search
7  grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=cv,
8  scoring='roc_auc')
9  # execute the grid search
10 grid_result = grid.fit(X_train, y_train)
11 # report the best configuration
12 print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
13 # report all configurations
14 means = grid_result.cv_results_['mean_test_score']
15 stds = grid_result.cv_results_['std_test_score']
16 params = grid_result.cv_results_['params']
17 for mean, stdev, param in zip(means, stds, params):
18     print("%f (%f) with: %r" % (mean, stdev, param))

```

```

1  # class imbalance with Gaussian
2
3  from numpy import sqrt
4  from numpy import argmax
5  from sklearn.datasets import make_classification
6  from sklearn.linear_model import LogisticRegression
7  from sklearn.model_selection import train_test_split
8  from sklearn.metrics import roc_curve
9  from matplotlib import pyplot
10
11
12 # fit a model
13 model = LogisticRegression(solver='lbfgs')
14 model.fit(X_train, y_train)
15 # predict probabilities
16 yhat = model.predict_proba(X_val)
17 # keep probabilities for the positive outcome only
18 yhat = yhat[:, 1]
19 # calculate roc curves
20 fpr, tpr, thresholds = roc_curve(y_val, yhat)

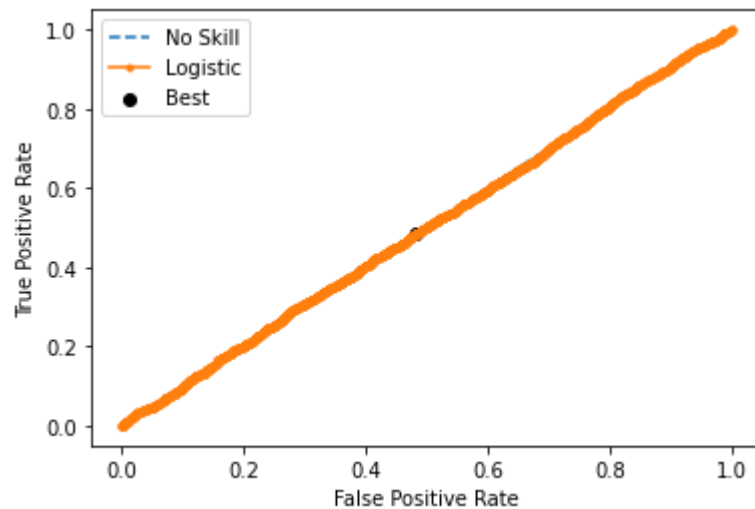
```

```

21 # calculate the g-mean for each threshold
22 gmeans = sqrt(tpr * (1-fpr))
23 # locate the index of the largest g-mean
24 ix = argmax(gmeans)
25 print('Best Threshold=%f, G-mean=%.3f' % (thresholds[ix], gmeans[ix]))
26 # plot the roc curve for the model
27 pyplot.plot([0,1], [0,1], linestyle='--', label='No Skill')
28 pyplot.plot(fpr, tpr, marker='.', label='Logistic')
29 pyplot.scatter(fpr[ix], tpr[ix], marker='o', color='black', label='Best')
30 # axis labels
31 pyplot.xlabel('False Positive Rate')
32 pyplot.ylabel('True Positive Rate')
33 pyplot.legend()
34 # show the plot
35 pyplot.show()

```

Best Threshold=0.252692, G-mean=0.501



```

1 # Class imbalance using SVM
2
3 from numpy import mean
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import RepeatedStratifiedKFold
7 from sklearn.calibration import CalibratedClassifierCV

```

```
7 from sklearn.calibration import CalibratedClassifierCV
8 from sklearn.svm import SVC

1 # define model
2 model = SVC(gamma='scale')
3 # wrap the model
4 calibrated = CalibratedClassifierCV(model, method='isotonic', cv=3)
5 # define evaluation procedure
6 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
7 # evaluate model
8 scores = cross_val_score(calibrated, X_train, y_train, scoring='roc_auc', cv=cv, n_jobs=-1)
9 # summarize performance
10 print('Mean ROC AUC: %.3f' % mean(scores))
```

/usr/local/lib/python3.7/dist-packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker stopped whi
"timeout or by a memory leak.", UserWarning

1 train_dataset

	City_Code	Region_Code	Accomodation_Type	Reco_Insurance_Type	Upper_Age	Lower_Age	Is_Spouse	Health Indicator	Holc
	0	22	3213.0	1	0	36.0	36.0	0	0

```

1 y = train_dataset.values
2 #train_dataset.drop(['ID', 'Response'], inplace=True, axis=1)
3
4 x = train_dataset.values
5

```

```

1 from lightgbm import LGBMClassifier
2
3 lgbm = LGBMClassifier(random_state=5)
4
5 lgbm.fit(X_train, y_train)
6
7 y_pred = lgbm.predict(X_val)


```

50881	0	729.0	1	0	25.0	25.0	0	0
--------------	---	-------	---	---	------	------	---	---

```

1 import sklearn.metrics as metrics
2 # calculate the fpr and tpr for all thresholds of the classification
3 probs = lgbm.predict_proba(X_val)
4 preds = probs[:,1]
5 fpr, tpr, threshold = metrics.roc_curve(y_val, preds)
6 roc_auc = metrics.auc(fpr, tpr)
7
8 # method I: plt
9 import matplotlib.pyplot as plt
10 plt.title('Receiver Operating Characteristic')
11 plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
12 plt.legend(loc = 'lower right')
13 plt.plot([0, 1], [0, 1], 'r--')
14 plt.xlim([0, 1])
15 plt.ylim([0, 1])
16 plt.ylabel('True Positive Rate')
17 plt.xlabel('False Positive Rate')
18 plt.show()

```

```
1 import sklearn.metrics as metrics
2 # calculate the fpr and tpr for all thresholds of the classification
3 probs = lgbm.predict_proba(test_dataset)
4 preds = probs[:,1]
5
```

```
1 ID = test_dataset['ID'].values
2
```

```
1 submission = pd.DataFrame({'ImageId':ID,
2                             'Response':preds,
3                             }).set_index('ImageId')
```

```
1 submission.head()
```

Response**ImageId****50883** 0.221873**50884** 0.368947

```
1 submission.to_csv('mnist_lgbm.csv', columns=['Response'])
```

```
50883 0.140000
```

These all observation.

1. Random Forest class_balanced will never impacted your model accuracy
2. class imbalance with Gaussian never gives us right accuracy
3. Weights Parameters doesn't work imbalance
4. Probabilistic methods also didn't solve the problem imbalance
5. Poor Performance on STOME
6. Poor Performance on SMOTE +SMOTETomek
7. No performance improvement over SMOTE over sampling
8. U never think of Performance if your using Undersampling
7. Lots of Books and Online class room talks about imbalance data. but it may work on small sample.
8. XG Boost works some better
9. LightGBM Worked very well.its master piece for class imbalance.
- 10.Even trained Model on Deep learning No improvement performance with help of

Pycaret having parameter class_imbalance is True, But No effect on IT.even pycaret works well because it handles XGBOOST, LIGHTGBM perfectly

AutoViml having class imbalance = SMOTE, I will never effect model ROC AUC score.

These all are my experience with analytics vidya competition.

ned Model on Deep learning No improvement performamce

Bottom line : Don't follow Bilindly on online platform posts. understand the core concept

Please find the experinments for the same.

This is formatted as c **bold text**ode

<https://neptune.ai/blog/lightgbm-parameters-guide>

1